

CS 149 // Lec 2 // A modern multicore processor

→ Instruction level parallelism

↳ superscalar execution (instruction level parallelism → ILP)

↳ automatically finds independent instructions

→ reading values from DRAM destroys it.

→ cache : data cache & instruction cache

↳ in the processor

↳ replacement policies include LRU

↳ hit
miss

↳ cold miss

↳ capacity miss

↳ conflict miss

↳ reduces latency of memory access

↳ also lower energy for lower level of cache

→ parallel processor

↳ duplicated with multiple instruction streams

↳ maybe 25% slower per "core" but overall

it is $\leftarrow 0.75 \times 2 = \underline{1.5x}$ faster!

↳ that is a speedup

→ without changing the instruction stream, the execution on the new processors will be 25% slower

- a thread is an instruction stream (each one)
- number of threads doesn't matter, we can get as many as we can in the hardware
- nvidia calls cores ← SMs (shared multiprocessors)
 - ↳ apple has heterogeneous cores (big + little)
- data parallelism
 - ↳ all threads do the same thing in terms of instructions
 - ↳ increase the number of ALUs and registers
- vector program ← using AVX instructions
 - ↳ uses the multitude of ALUs
 - ↳ can use a single instruction stream
 - ↳ and then multiple processors can have multiple instruction streams.
- SIMD ← single instruction, multiple data
- ALUs are very cheap, FPUs are not as much
- what about conditionals?
 - ↳ mark outputs of unneeded vector bits
 - ↳ this is wasteful, we could get up to $1/8^{\text{th}}$ performance
 - ↳ instruction stream coherence $1/8^{\text{th}}$ divergent

